

```
In [1]: import os
print(os.getcwd())
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV
%matplotlib inline
```

C:\Users\riyav

```
In [2]: df = pd.read_csv('car_evaluation.csv', header = None)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	0	1	2	3	4	5	6
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [4]: col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
df.columns = col_names
col_names
```

```
Out[4]: ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

```
In [5]: df.head()
```

```
Out[5]:
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   object
1   maint       1728 non-null   object
2   doors       1728 non-null   object
3   persons     1728 non-null   object
4   lug_boot    1728 non-null   object
5   safety      1728 non-null   object
6   class       1728 non-null   object
```

dtypes: object(7)  
memory usage: 94.6+ KB

```
In [7]: for i in col_names:
        print(df[i].value_counts())
```

```
vhigh    432
high     432
med      432
low      432
Name: buying, dtype: int64
vhigh    432
high     432
med      432
low      432
Name: maint, dtype: int64
2        432
3        432
4        432
5more    432
Name: doors, dtype: int64
2        576
4        576
more     576
Name: persons, dtype: int64
small    576
med      576
big      576
Name: lug_boot, dtype: int64
low      576
med      576
high     576
Name: safety, dtype: int64
unacc   1210
acc      384
good      69
vgood     65
Name: class, dtype: int64
```

```
In [8]: df.shape
```

```
Out[8]: (1728, 7)
```

```
In [9]: X = df.drop(['class'],axis = 1)
        y = df['class']
```

```
In [10]: from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=42)
```

```
In [11]: from sklearn.preprocessing import OrdinalEncoder
         enc = OrdinalEncoder()
         X_train = enc.fit_transform(X_train)
         X_test = enc.transform(X_test)
```

## Gini index as criterion

```
In [12]: from sklearn.tree import DecisionTreeClassifier
```

```
In [13]: clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=42)
         clf_gini.fit(X_train, y_train)
```

```
Out[13]: DecisionTreeClassifier(max_depth=3, random_state=42)
```

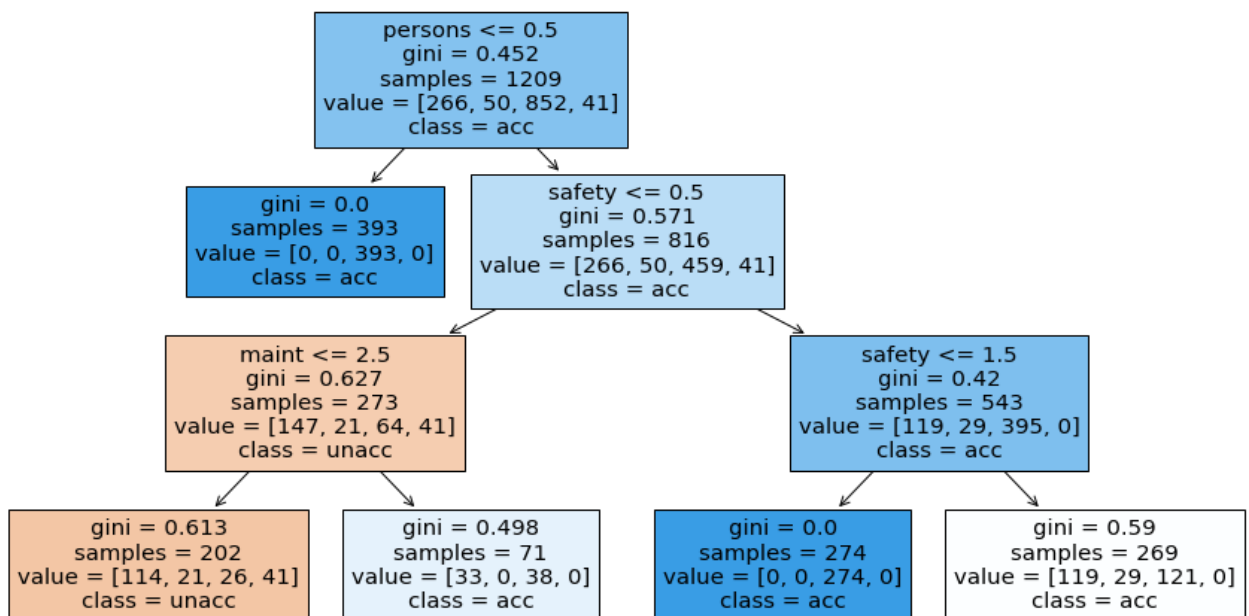
```
In [14]: y_pred = clf_gini.predict(X_test)
```

```
In [15]: from sklearn.metrics import accuracy_score
```

```
print(f'Model with gini index gives an accuracy of: {accuracy_score(y_true=y_test, y_pr
```

Model with gini index gives an accuracy of: 0.7572254335260116

```
In [16]: from sklearn import tree
plt.figure(figsize=(15,8))
tree.plot_tree(clf_gini,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safet
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```



```
In [17]: # Check for underfitting
```

```
print(f'Training set score: {clf_gini.score(X_train,y_train)}')
print(f'Test set score: {clf_gini.score(X_test,y_test)}')
```

Training set score: 0.7775020678246485  
Test set score: 0.7572254335260116

## Entropy as criterion

```
In [18]: clf_entropy = DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
clf_entropy.fit(X_train, y_train)
```

```
Out[18]: DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
```

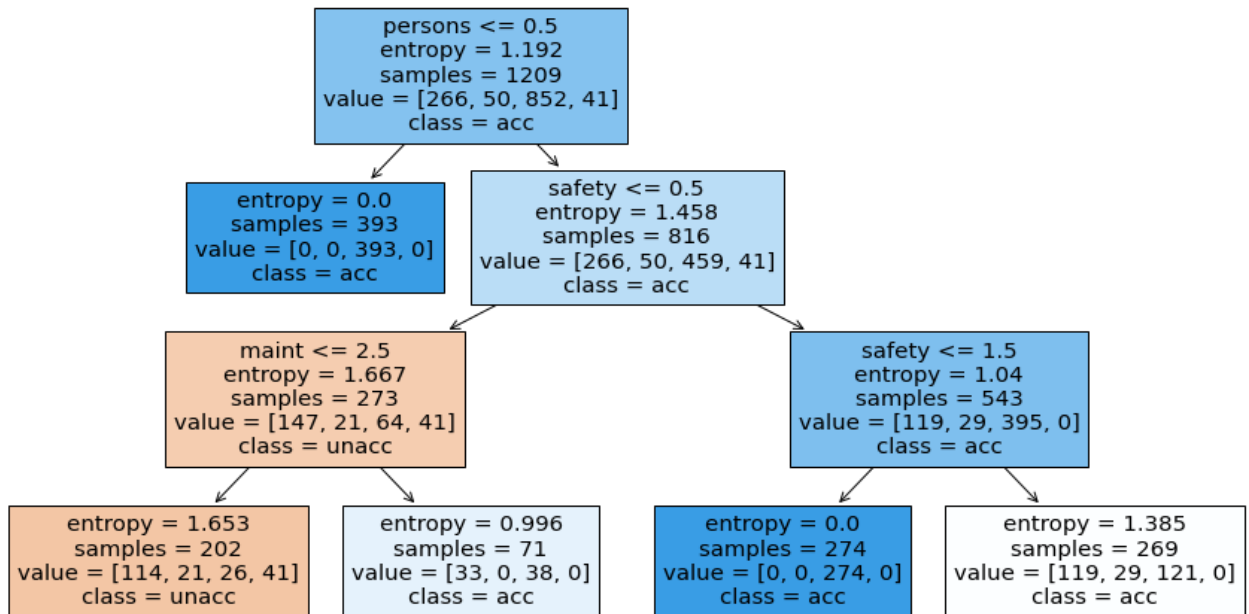
```
In [19]: y_pred = clf_entropy.predict(X_test)
```

```
In [20]: from sklearn.metrics import accuracy_score
```

```
print(f'Model with gini index gives an accuracy of: {accuracy_score(y_test, y_pred)}')
```

Model with gini index gives an accuracy of: 0.7572254335260116

```
In [21]: plt.figure(figsize=(15,8))
         tree.plot_tree(clf_entropy,
                       feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safet',
                                       'class_names= list(set(y_train)),
                                       filled = True)
         plt.show()
```



```
In [22]: # Check for underfitting

         print(f'Training set score: {clf_entropy.score(X_train,y_train)}')
         print(f'Test set score: {clf_entropy.score(X_test,y_test)}')
```

Training set score: 0.7775020678246485

Test set score: 0.7572254335260116

```
In [23]: from sklearn.metrics import confusion_matrix, classification_report
         cm = confusion_matrix(y_test, y_pred)
```

```
In [24]: print(cm)
```

```
[[ 44  0  74  0]
 [  9  0  10  0]
 [  9  0 349  0]
 [ 24  0  0  0]]
```

```
In [25]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
acc	0.51	0.37	0.43	118
good	0.00	0.00	0.00	19
unacc	0.81	0.97	0.88	358
vgood	0.00	0.00	0.00	24
accuracy			0.76	519
macro avg	0.33	0.34	0.33	519

weighted avg	0.67	0.76	0.71	519
--------------	------	------	------	-----

```
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: U
ndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: U
ndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\PRATYUSH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1248: U
ndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in lab
els with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## Cross Validation

```
In [26]: params_grid = {
        'criterion':['gini','entropy'],
        'max_depth':[3,4,5,6,7,8,9,10]
        }
```

```
In [27]: decision_tree = DecisionTreeClassifier()
        decision_tree.fit(X_train,y_train)
```

```
Out[27]: DecisionTreeClassifier()
```

```
In [28]: dt_validated = GridSearchCV(estimator=decision_tree, param_grid=params_grid, scoring='ac
```

```
In [29]: %%time
        dt_validated.fit(X_train,y_train)
```

Wall time: 840 ms

```
Out[29]: GridSearchCV(cv=20, estimator=DecisionTreeClassifier(),
        param_grid={'criterion': ['gini', 'entropy'],
        'max_depth': [3, 4, 5, 6, 7, 8, 9, 10]},
        scoring='accuracy')
```

```
In [30]: print(f'Best parameters for decison tree classifier after CV -> {dt_validated.best_para
        print(f'Best score on decision tree classifier after CV -> {dt_validated.best_score_}')
```

Best parameters for decison tree classifier after CV -> {'criterion': 'entropy', 'max\_de  
pth': 10}

Best score on decision tree classifier after CV -> 0.979330601092896

```
In [31]: print(f'Score on train set of DT classifier before CV -> {decision_tree.score(X_train,
        print(f'Score on test set of DT classifier before CV -> {decision_tree.score(X_test, y_
        print(f'Score on train set of DT classifier after CV -> {dt_validated.score(X_train, y_
        print(f'Score on test set of DT classifier after CV -> {dt_validated.score(X_test, y_te
```

Score on train set of DT classifier before CV -> 1.0

Score on test set of DT classifier before CV -> 0.9653179190751445

Score on train set of DT classifier after CV -> 0.9925558312655087

Score on test set of DT classifier after CV -> 0.9595375722543352

```
In [32]: print('Classification report on train set')
        print(classification_report(y_true=y_train, y_pred=dt_validated.predict(X_train)))
```

Classification report on train set

precision	recall	f1-score	support
-----------	--------	----------	---------

acc	0.98	0.99	0.99	266
good	0.98	0.98	0.98	50
unacc	1.00	0.99	1.00	852
vgood	0.98	1.00	0.99	41
accuracy			0.99	1209
macro avg	0.98	0.99	0.99	1209
weighted avg	0.99	0.99	0.99	1209

```
In [33]: print('Classification report on test set')
print(classification_report(y_true=y_test, y_pred=dt_validated.predict(X_test)))
```

Classification report on test set

	precision	recall	f1-score	support
acc	0.92	0.92	0.92	118
good	0.71	0.89	0.79	19
unacc	0.99	0.98	0.99	358
vgood	0.88	0.88	0.88	24
accuracy			0.96	519
macro avg	0.88	0.92	0.89	519
weighted avg	0.96	0.96	0.96	519

```
In [34]: print('Confusion matrix on train set')
print(confusion_matrix(y_true=y_train, y_pred=dt_validated.predict(X_train)))
```

Confusion matrix on train set

```
[[263  1  2  0]
 [  0 49  0  1]
 [  5  0 847  0]
 [  0  0  0 41]]
```

```
In [35]: print('Confusion matrix on test set')
print(confusion_matrix(y_true=y_test, y_pred=dt_validated.predict(X_test)))
```

Confusion matrix on test set

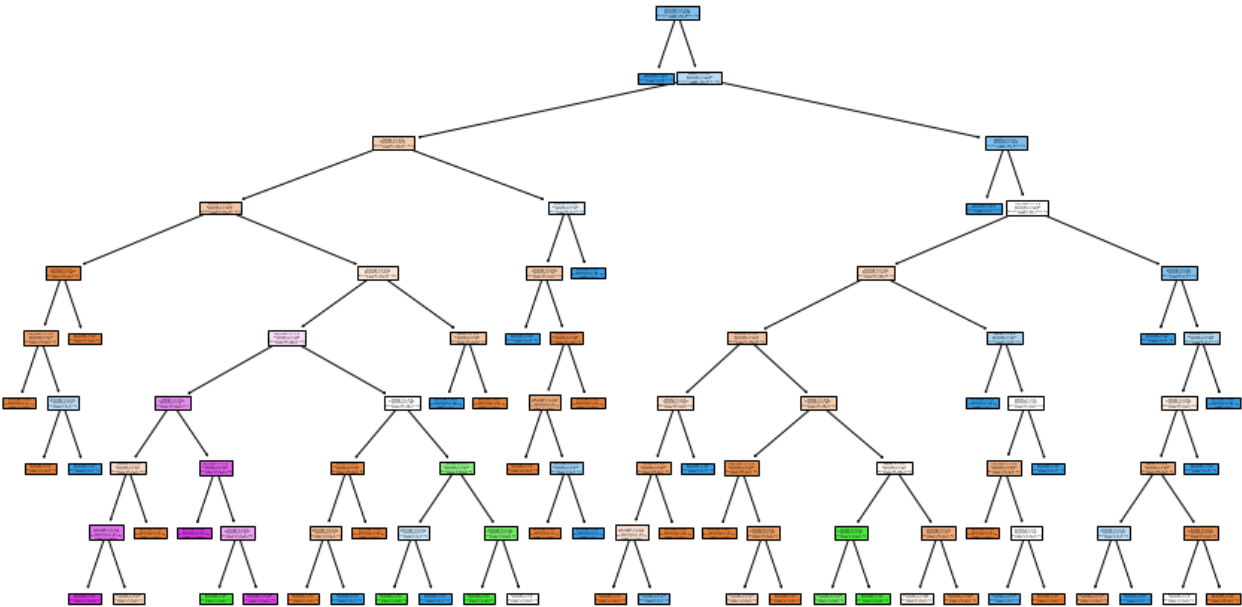
```
[[108  7  2  1]
 [  0 17  0  2]
 [  6  0 352  0]
 [  3  0  0 21]]
```

```
In [39]: best_dt = DecisionTreeClassifier(criterion='entropy',max_depth=9)
```

```
In [40]: best_dt.fit(X_train,y_train)
```

```
Out[40]: DecisionTreeClassifier(criterion='entropy', max_depth=9)
```

```
In [41]: plt.figure(figsize=(15,8))
tree.plot_tree(best_dt,
                feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safet',
                               'class_names= list(set(y_train)),
                               filled = True)
plt.show()
```



In [ ]: